



RAHUL K

Azure Data Engineer
TCS | 3.5+ Years

CONTACT

+91 8156886640
rahulraghavankky@gmail.com
linkedin.com/in/rahul-k-detcs
Bengaluru, Karnataka

CORE SKILLS & TECHNOLOGIES

- Azure Data Factory (ADF)
- Azure Data Lake Gen2 (ADLS)
- Azure Blob Storage
- Azure Key Vault
- Microsoft Fabric
- Azure Databricks / PySpark
- Delta Lake / Lakehouse
- Medallion Architecture
- Metadata-Driven Pipelines
- SQL Server & Oracle 13c
- Python & Spark SQL
- ETL / ELT Processes
- CI/CD – Azure DevOps / Git
- Dimensional Modelling

CERTIFICATIONS

Databricks Certified

Assoc. Data Engineer

AWS Certified

Developer Associate

(DVA-C02)

Microsoft Certified

Azure Data Fundamentals

(DP- 900)

EDUCATION

B.Tech — Computer Science

LBS College of Engineering

7.7cgsa(2018-22)

Durga Higher Secondary School

(PCM(CS)) – 95%

(2016-18)

Government Higher Secondary

School

(PCMS) – 98%

(2016)

LANGUAGES

- English
- Malayalam
- Hindi

PROFESSIONAL SUMMARY

Certified Azure Data Engineer with 3.5+ years of experience delivering enterprise-grade data pipelines and lakehouse architectures across Aviation and BFSI domains. Expert in ADF, PySpark, Delta Lake, Metadata-Driven pipeline design, Microsoft Fabric, and CI/CD via Azure DevOps. Certified in Databricks, Microsoft Fabric, AWS Developer Associate, and DP-900 — skilled at building scalable, secure, and cost-efficient data platforms.

WORK EXPERIENCE 1

Azure Data Engineer

Tata Consultancy Services (TCS)

Jun 2023 – Present

Domain: BFSI | Global Banking Client | Azure Data Platform Modernization

- Architected 20+ parameterized ADF pipelines ingesting data from SQL Server, Oracle and REST APIs into Azure Data Lake Gen2 — processing 40M+ records daily across banking and insurance workloads
- Built PySpark transformation layers on Azure Databricks following Medallion Architecture (Bronze → Silver → Gold), improving data quality and query performance by 50%
- Implemented Delta Lake with SCD Type 2 logic for ACID-compliant historical tracking of customer and transaction data across 10+ dimension tables
- Integrated Azure Key Vault across 30+ pipeline components, eliminating hardcoded credentials and reducing exposure risks by 100%
- Established CI/CD via Azure DevOps and Git — reducing release cycle from 2 weeks to 3 days across Dev / UAT / Production environments
- Delivered Microsoft Fabric Lakehouse and Dataflows Gen2 for self-service analytics, cutting business dependency on engineering team by 40%
- Migrated legacy Oracle ETL to ADF + Databricks — reducing infrastructure costs by 30% and pipeline run time by 45%

WORK EXPERIENCE 2

Azure Data Engineer

Lenok Solutions

Aug 2022 – Mar 2023

Domain: Aviation | Airline Operations, Passenger Analytics & MRO Intelligence

- Designed a Metadata-Driven Incremental ADF Pipeline — watermark-based, zero-hardcoded framework dynamically processing 25+ aviation SQL tables (schedules, booking, MRO) into ADLS Gen2, cutting pipeline runtime by 60% vs full loads
- Built Lookup + ForEach + Conditional ADF logic via metadata control table (TableName, WatermarkColumn, LastLoadValue, IsActive) — enabling seamless new source onboarding without any pipeline changes
- Implemented Self-Hosted Integration Runtime (SHIR) for secure encrypted data movement from on-prem airline databases to Azure cloud
- Developed PySpark jobs on Azure Databricks processing 10M+ daily flight records for route profitability, seat utilization and on-time performance analytics
- Integrated Azure Key Vault for all credential management and implemented CI/CD via Azure DevOps with ARM templates across Dev / Test / UAT / Prod
- Optimized Blob Storage tiers (Hot/Cool) and lifecycle policies — cutting cloud storage costs by 40% without impacting data accessibility